



Scalable Machine Learning: Distributed Training Methods for Large Models and Datasets

Presentation by:

- Mohamed STIFI

10/01/2025

Overview

01 distributed training

02 Experimental Environment

03 Experimental Results

04 Results Analysis



01


distributed training

01

What is distributed training

Vertical scaling

No code change




1x Server
8GB RAM
4GB GPU Memory

1x Server
64GB RAM
32GB GPU Memory

Horizontal scaling

Minimal code change (thanks to PyTorch)




1x Server
8GB RAM
4GB GPU Memory

4x Servers
8GB RAM
4GB GPU Memory (x2)


01

Data Parallelism vs Model Parallelism

Data parallelism



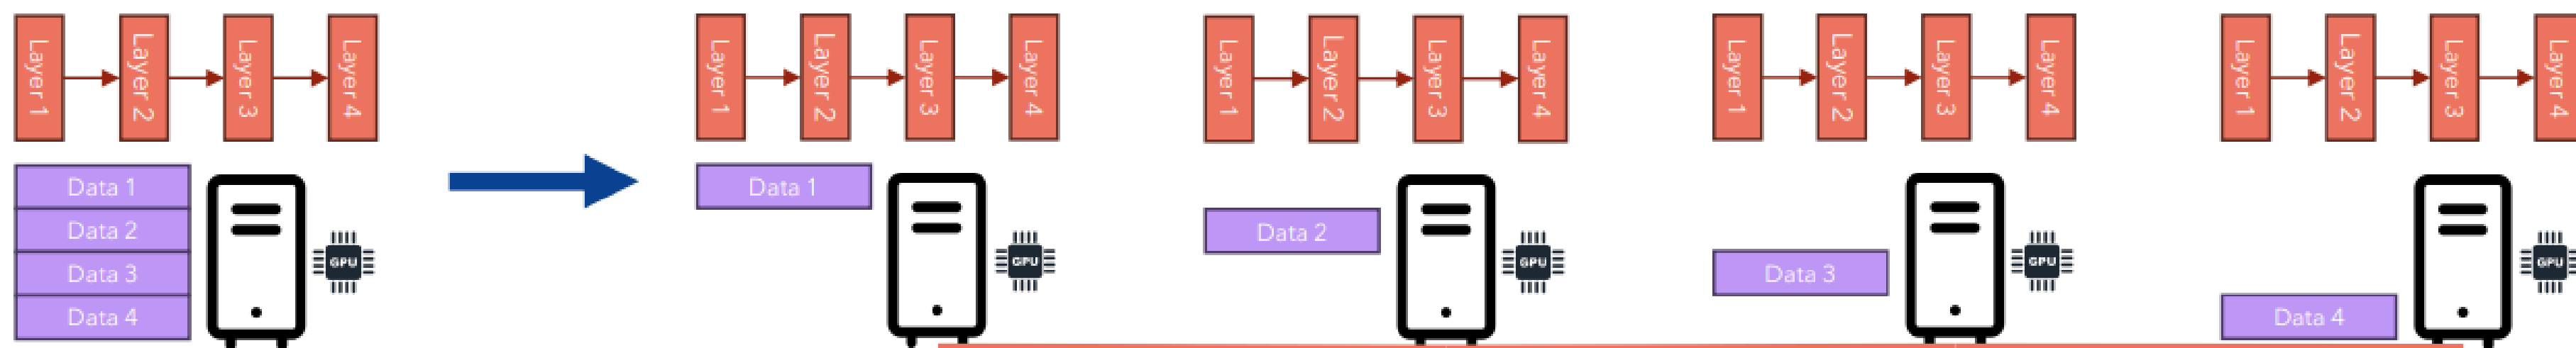
Model parallelism



01



Data Parallelism

If the model can fit within a single GPU, then we can distribute the training on multiple servers (each containing one or multiple GPUs), with each GPU processing a subset of the entire dataset in parallel and synchronizing the gradients during backpropagation. This option is known as Data Parallelism.






01

Types of Data Parallelism



Single-Server, Multi-GPU

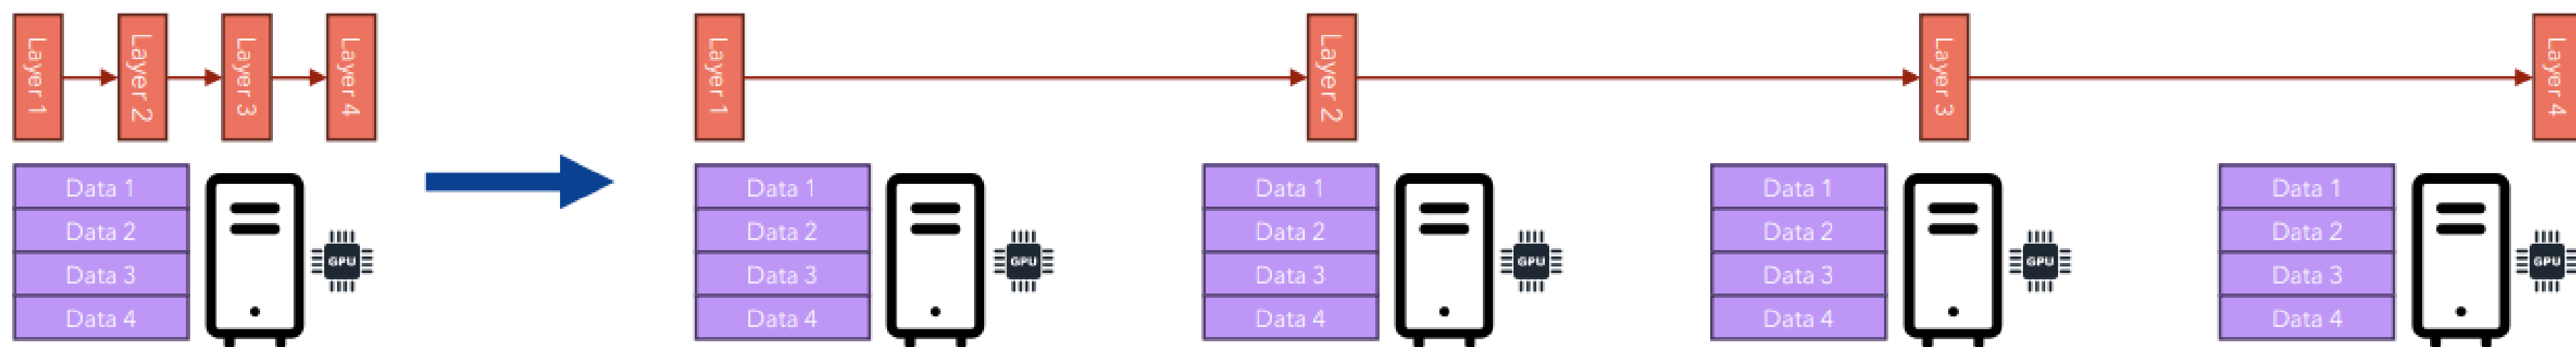


Multi-Server, Multi-GPU

01


Model Parallelism

If the model cannot fit within a single GPU, then we need to “break” the model into smaller layers and let each GPU process a part of the forward/backward step during gradient descent. This option is known as Model Parallelism.



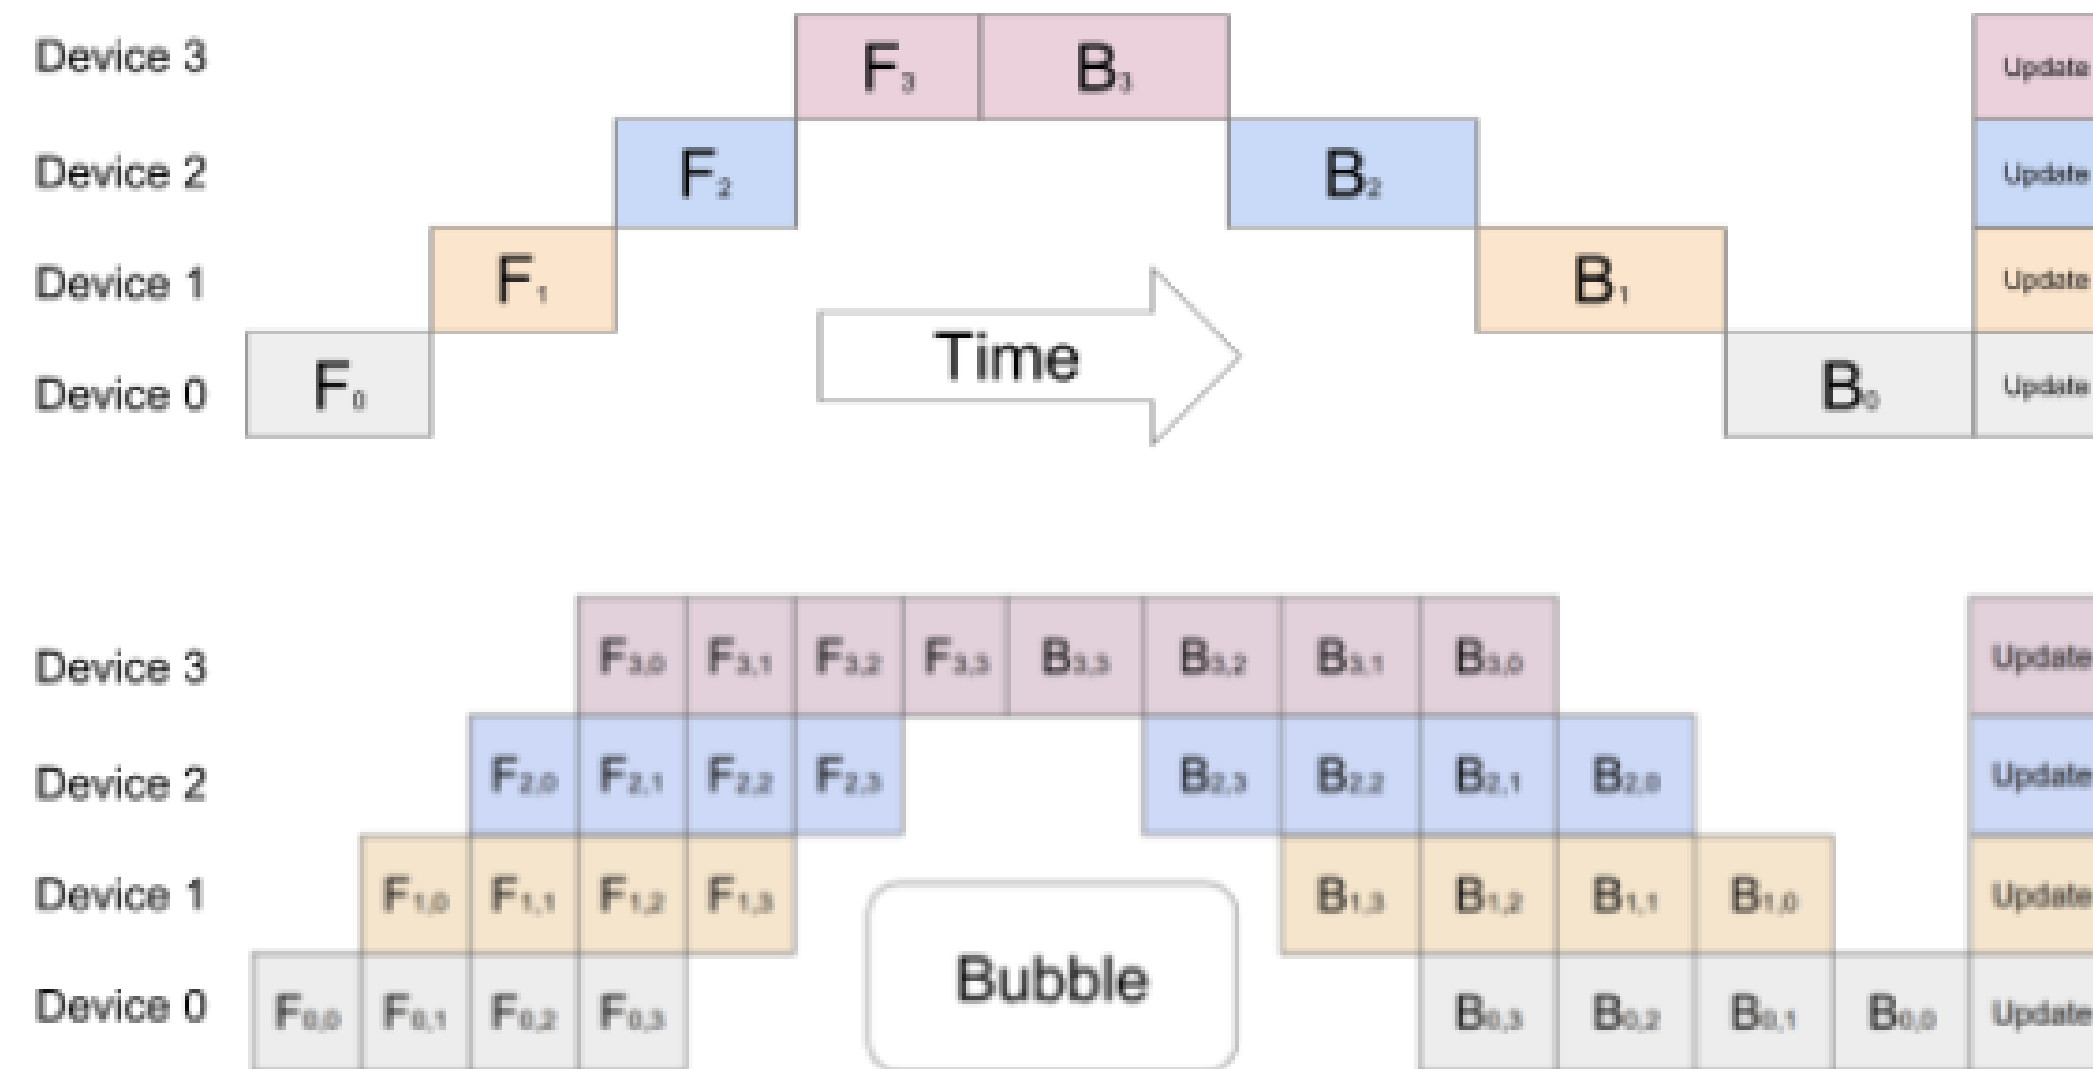
01

Types of Model Parallelism



Layer-wise Parallelism

Types of Model Parallelism



Pipeline Parallelism

01

Algorithm

Algorithm 1 Algorithm for Distributed Training in Machine Learning

Require: Number of GPUs *world_size*, Dataset *D*, Number of Epochs *num_epochs*, Learning Rate *LR*

Ensure: Trained Model *M*



- 1: **Set Up the Environment :** ▷ Import libraries and define global variables
 - 2: **Initialize the Distributed Data Parallel (DDP) Environment :**
 - 3: Define and execute functions :
 - 4: `setup(rank, world_size)`: Initialize process group
 - 5: `cleanup()`: Destroy process group
 - 6: **Define a Model :**
 - 7: Implement architecture with `Model(nn.Module)`
 - 8: For model parallelism, manually assign layers to devices
 - 9: `create_model()`: Return an instance of the model
 - 10: **Create a Dataloader :**
 - 11: Partition dataset *D* across GPUs using `DistributedSampler`
 - 12: `create_dataloader(rank, world_size)`:
 - 13: Partition *D*, create mini-batches, and return dataloader instances
 - 14: **Implement the Training Loop :**
 - 15: Define helper functions :
 - 16: `train(model, iterator, optimizer, criterion, rank)`: Single training step
 - 17: `evaluate(model, iterator, criterion, rank)`: Single evaluation step
 - 18: Define main training function `main_train(rank, world_size)`:
 - 19: a. Setup distributed process groups with `setup(rank, world_size)`
 - 20: b. Create model and dataloaders
 - 21: c. Wrap model with `DistributedDataParallel (DDP)`
 - 22: d. Define loss criterion and optimizer
 - 23: e. Train for *num_epochs*, compute metrics
 - 24: f. Evaluate on test set after training
 - 25: g. Cleanup environment with `cleanup()`
 - 26: **Main Execution :**
 - 27: Define execution function :
 - 28: Set *world_size* (number of GPUs)
 - 29: Use `multiprocessing.spawn()` to start distributed processes
 - 30: Ensure algorithm supports model and data parallelism
-

02

Experimental Environment

02

Experimental Environment



kaggle

Experimental Environment

Models:

- *VGG11*
- *MLP*
- *NBoW*
- *CNN*
- *Seq2Seq*
- *ConvAutoEncoder*

metrics:

- *Training Time*
- *Evaluation Time*
- *Loss*
- *Accuracy*


03

Experimental Results


VGG11

Training and Validation Accuracy

VGG11 - Training Accuracy

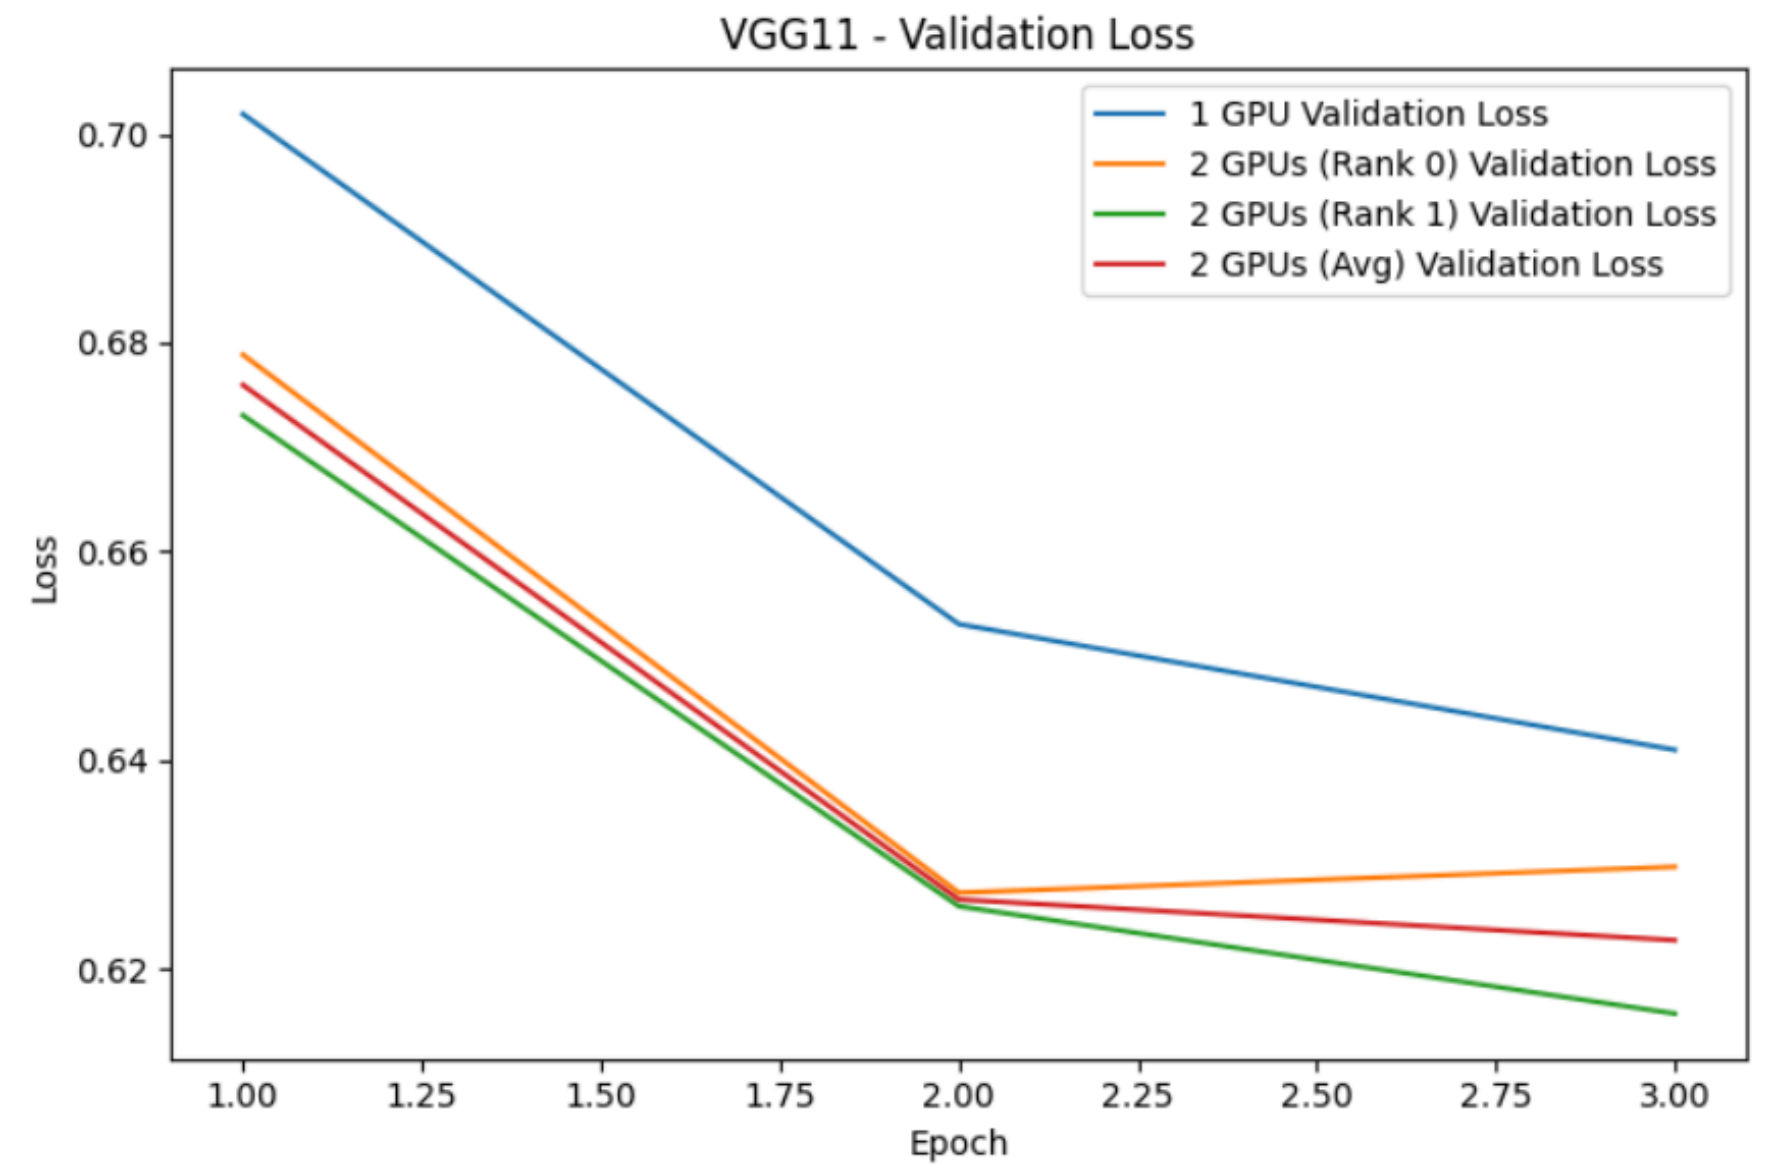
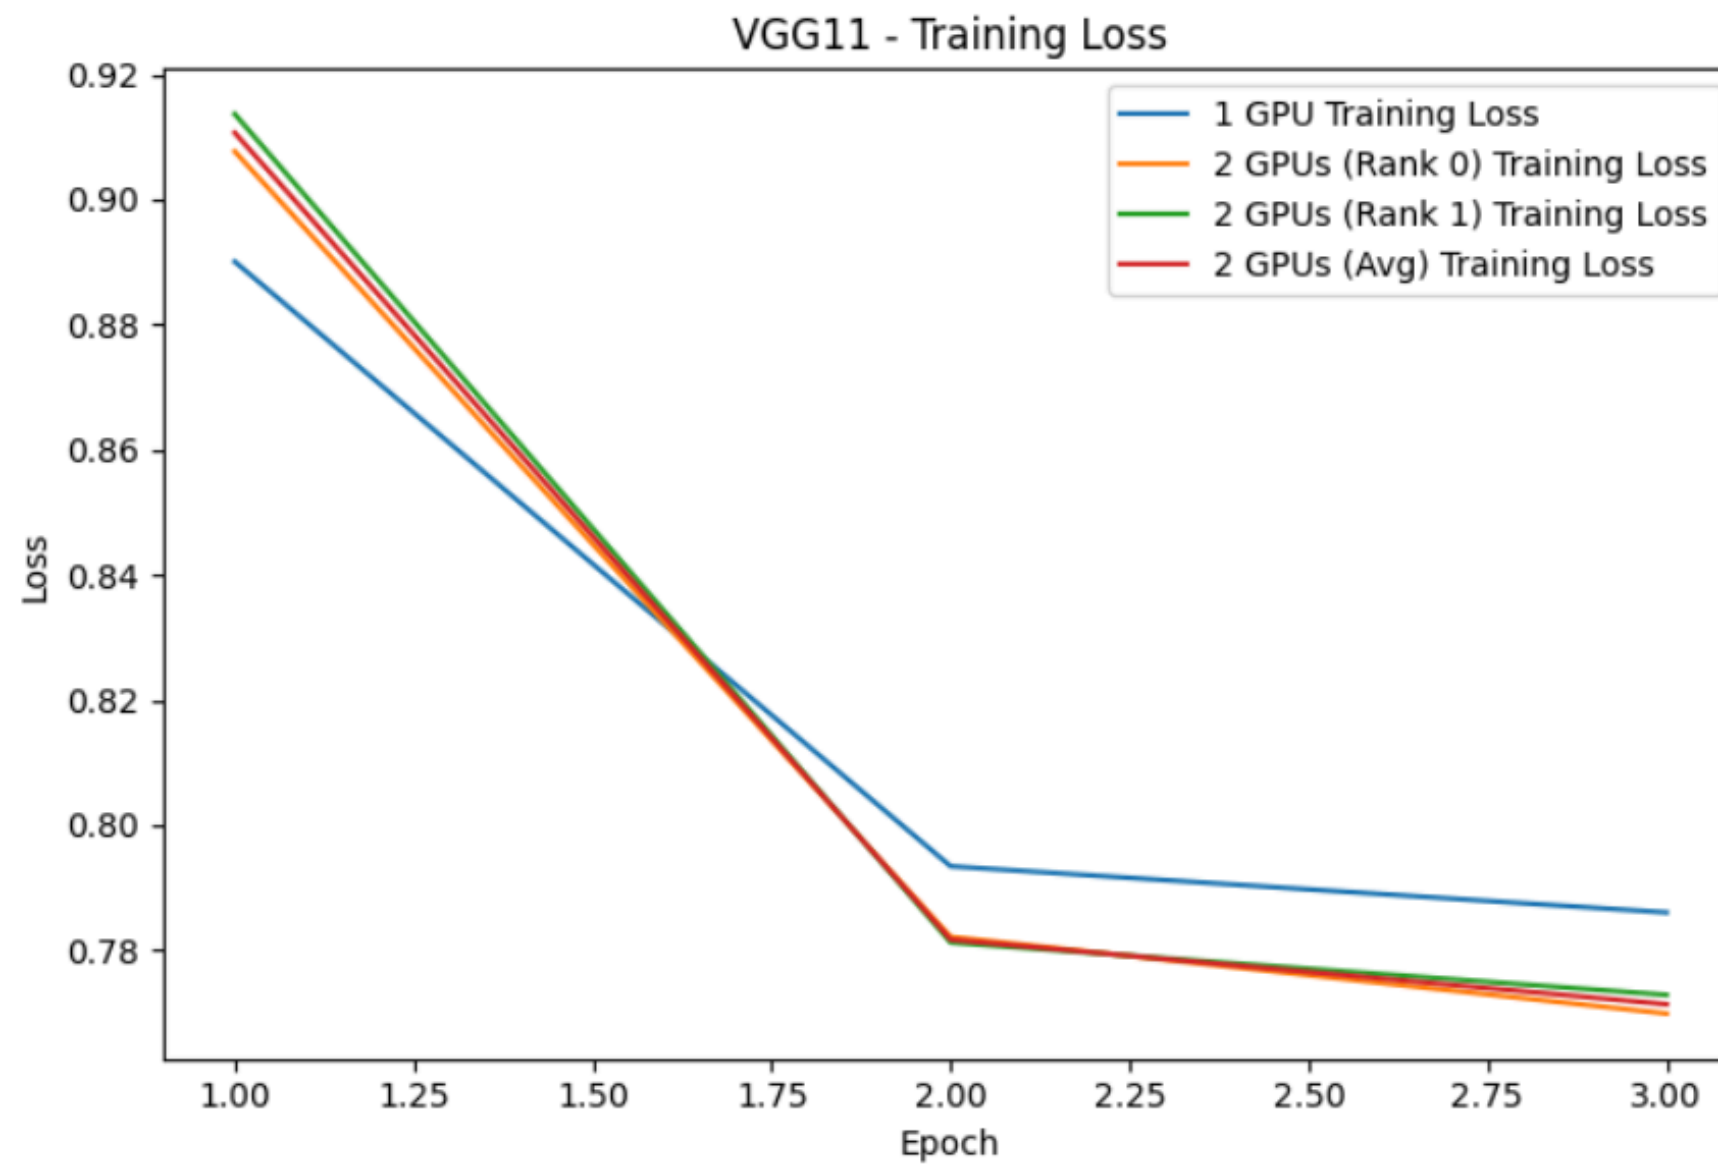


VGG11 - Validation Accuracy



VGG11

Training and Validation Loss

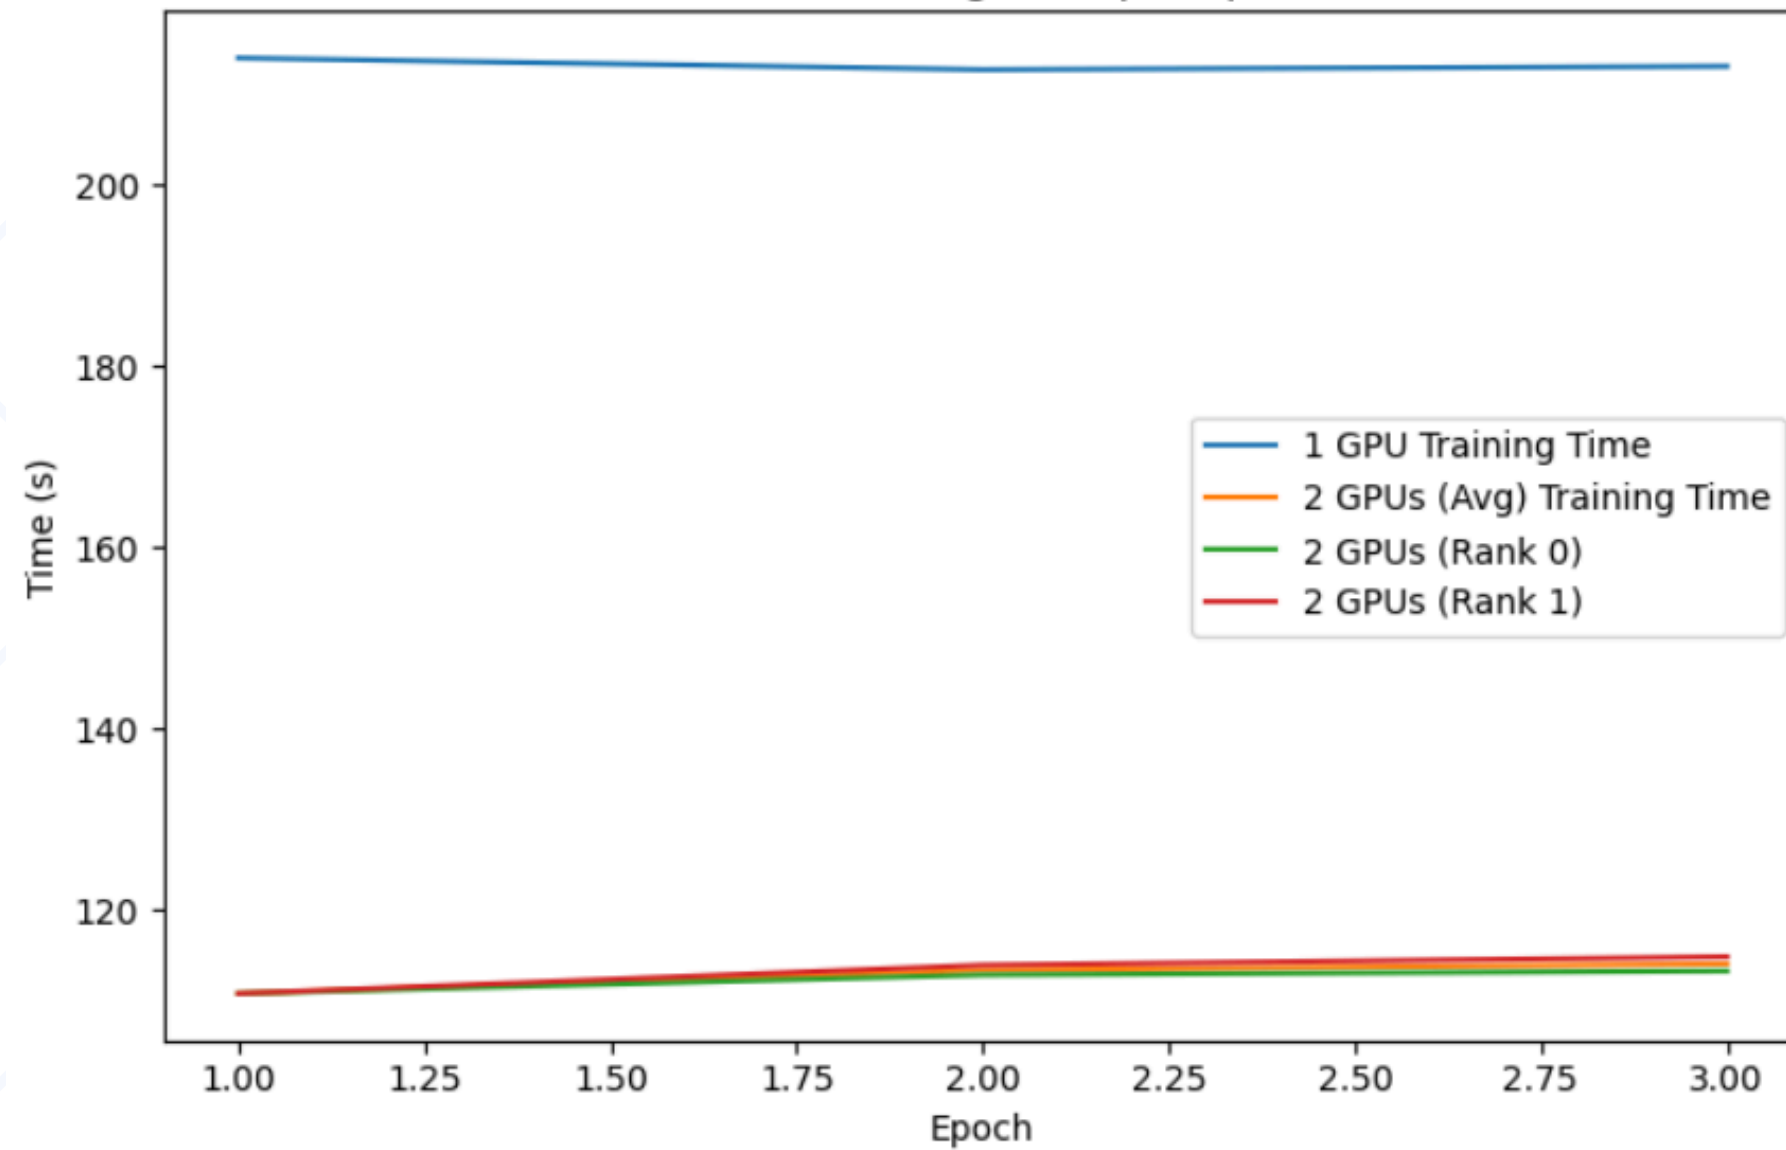


03

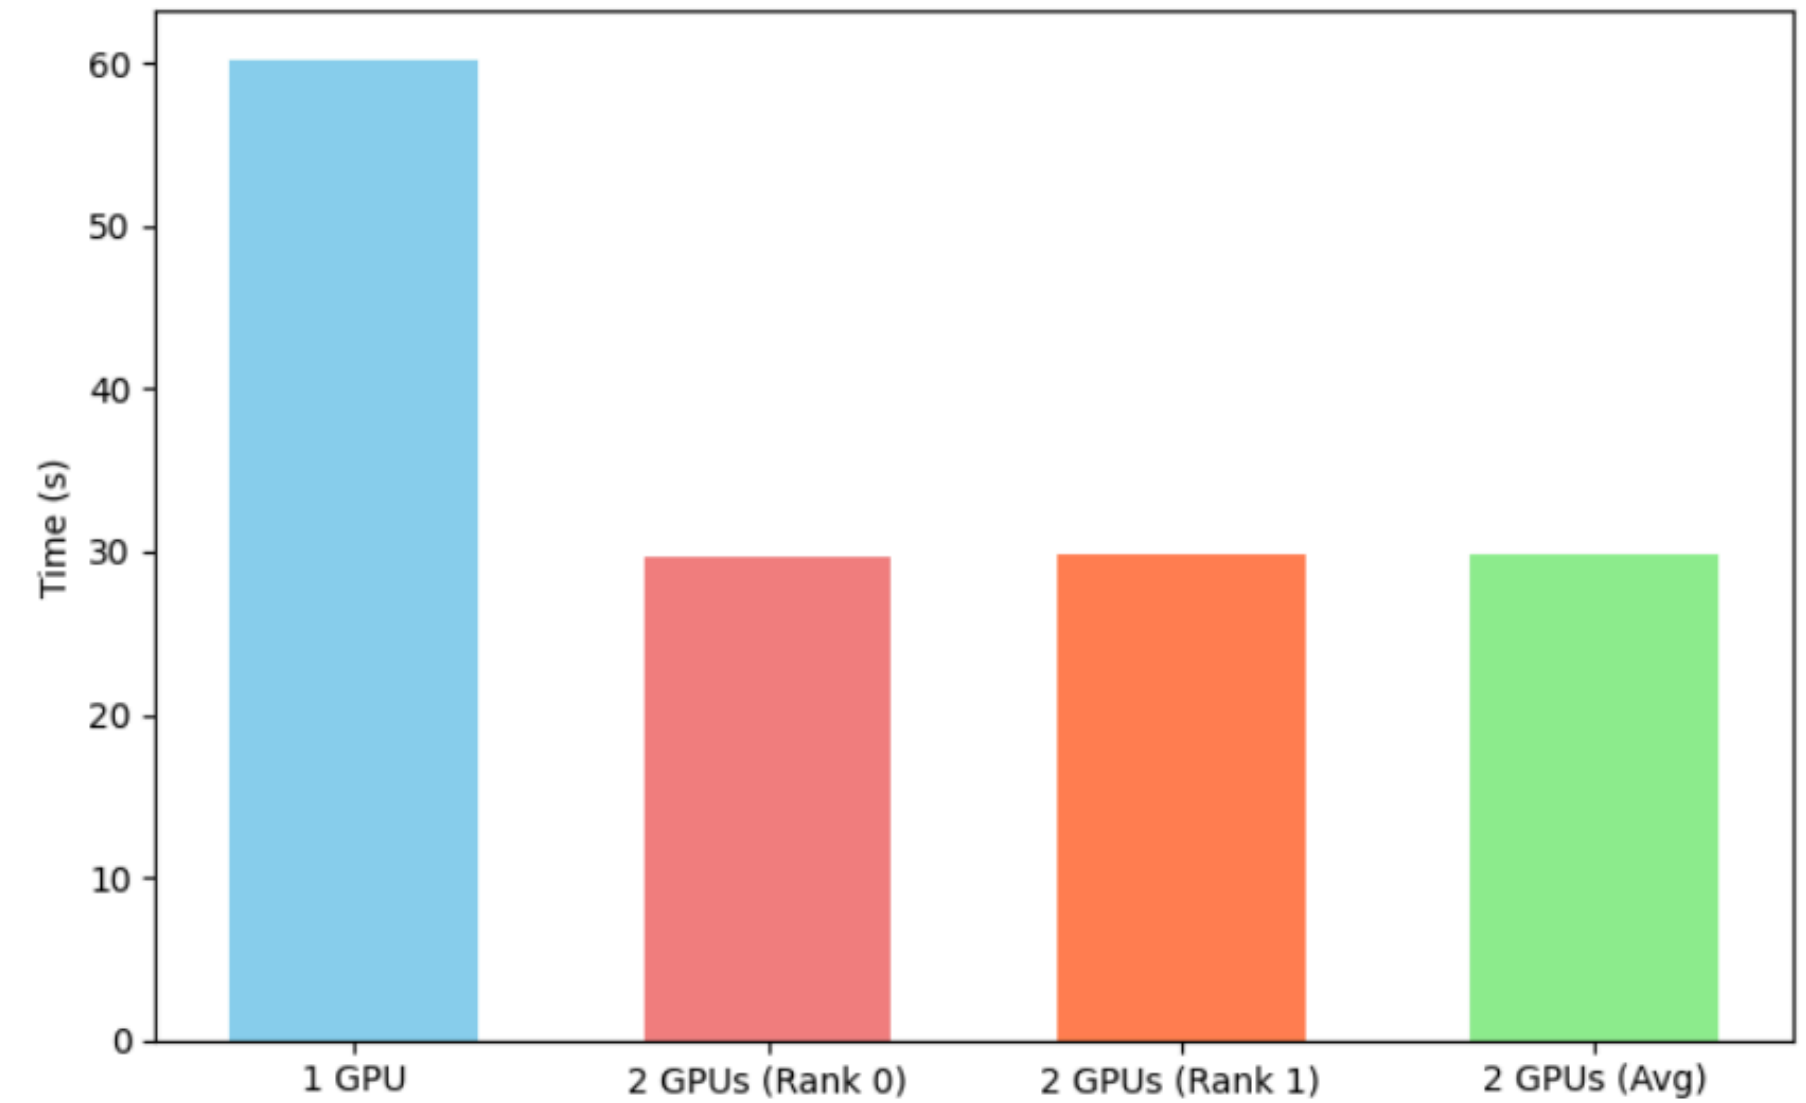
VGG11

Training and Validation Time

VGG11 - Training Time per Epoch



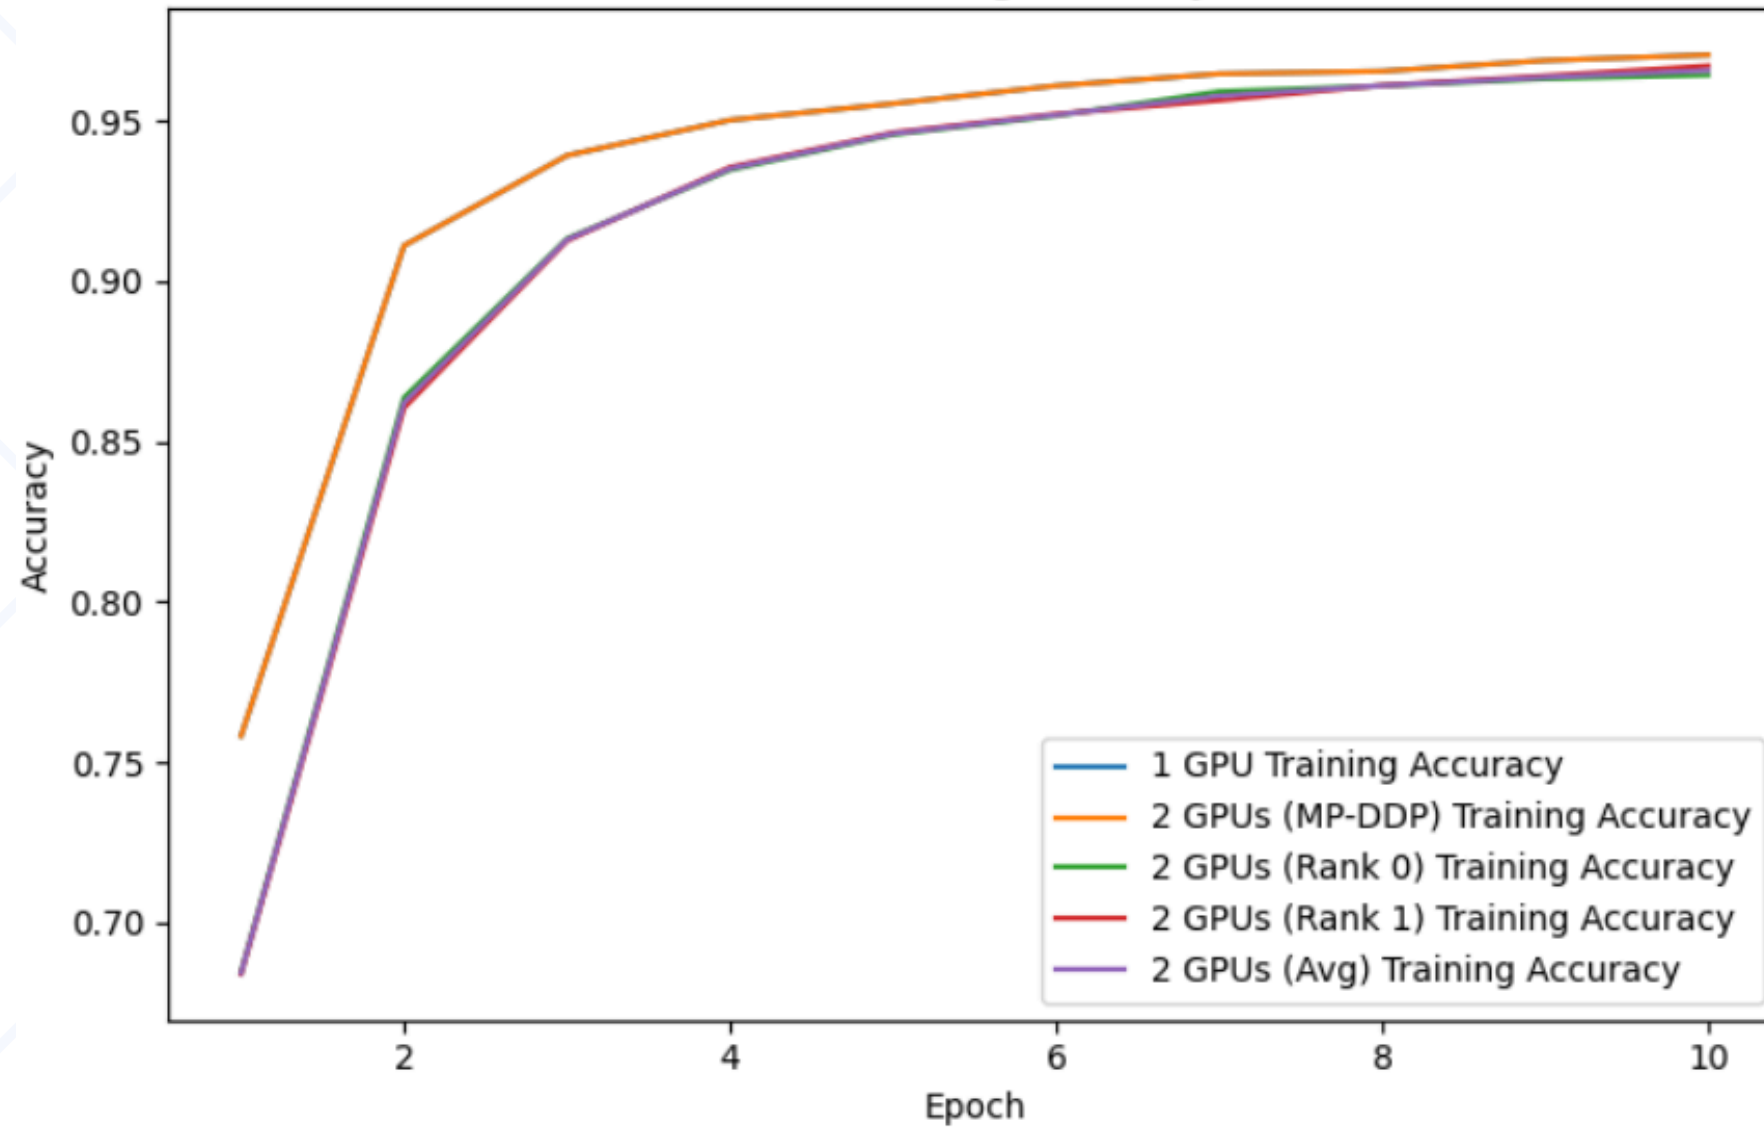
VGG11 - Total Validation Time



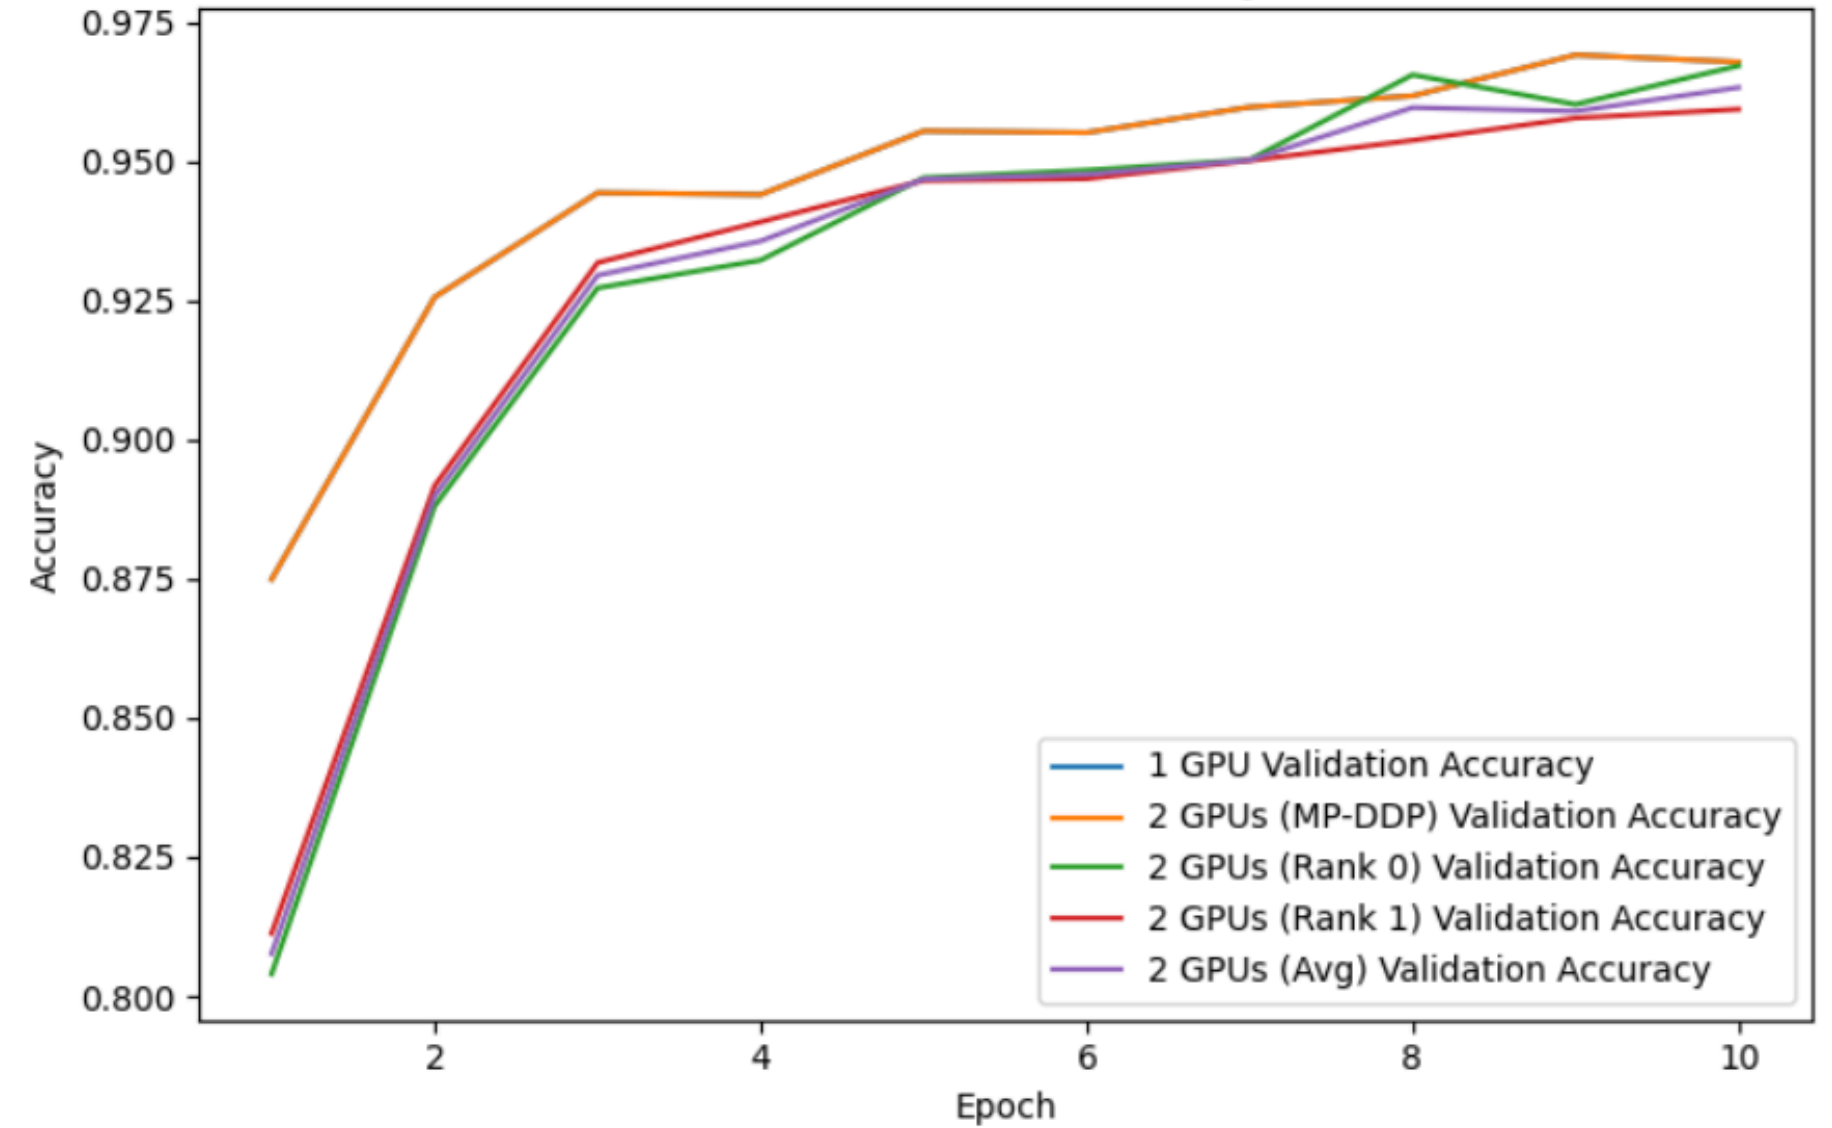
MLP

Training and Validation Accuracy

MLP - Training Accuracy



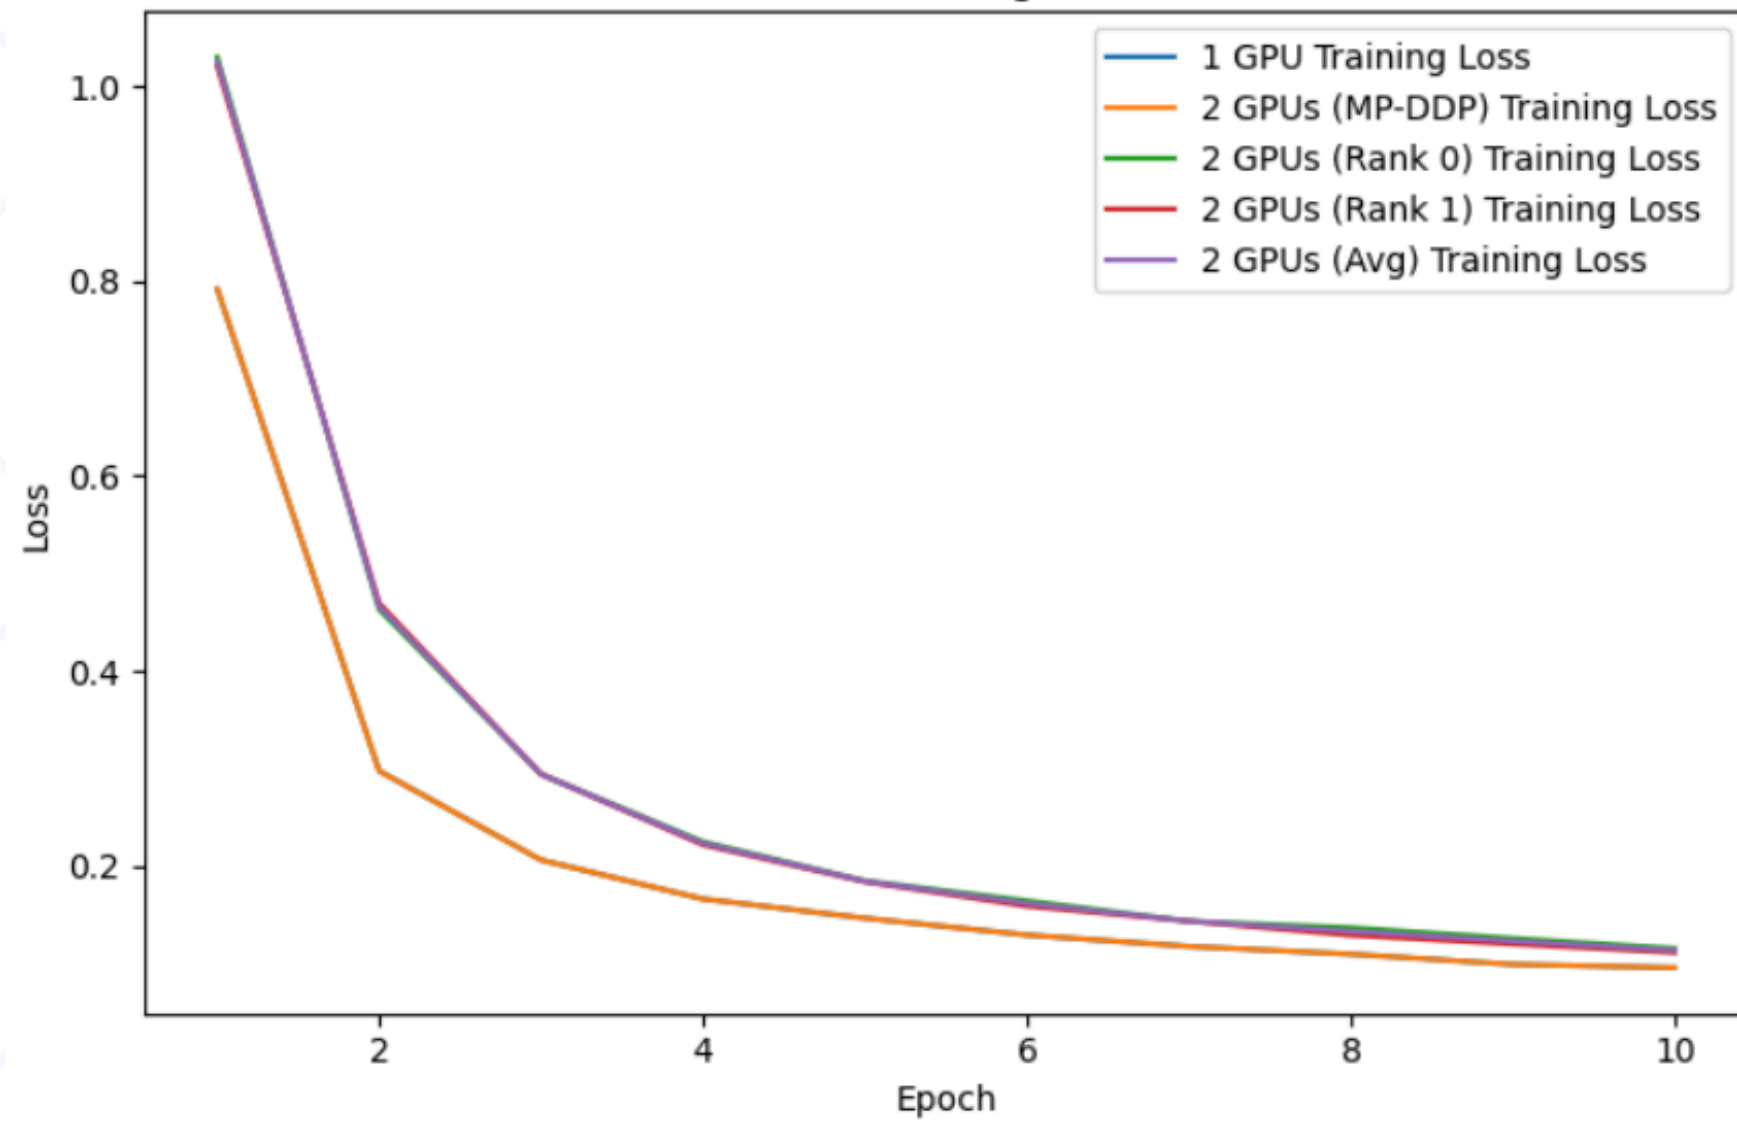
MLP - Validation Accuracy



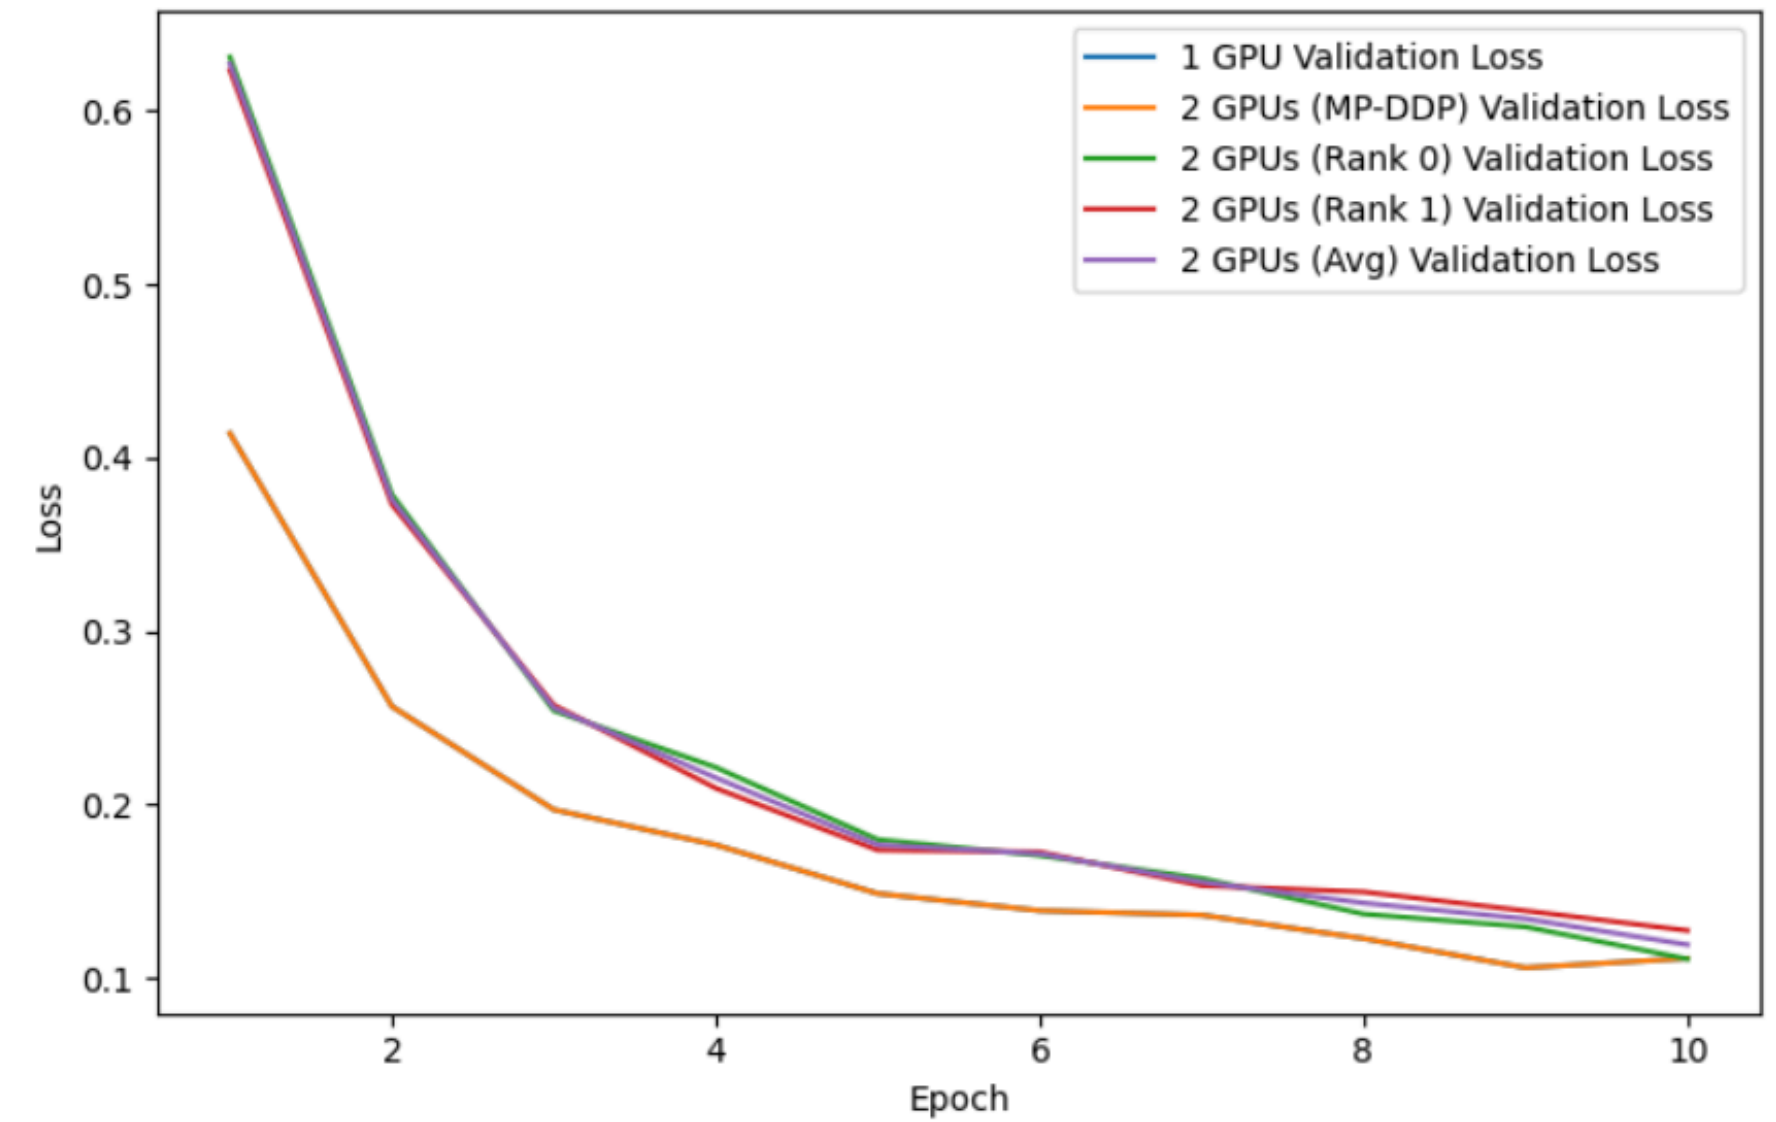
MLP

Training and Validation Loss

MLP - Training Loss



MLP - Validation Loss

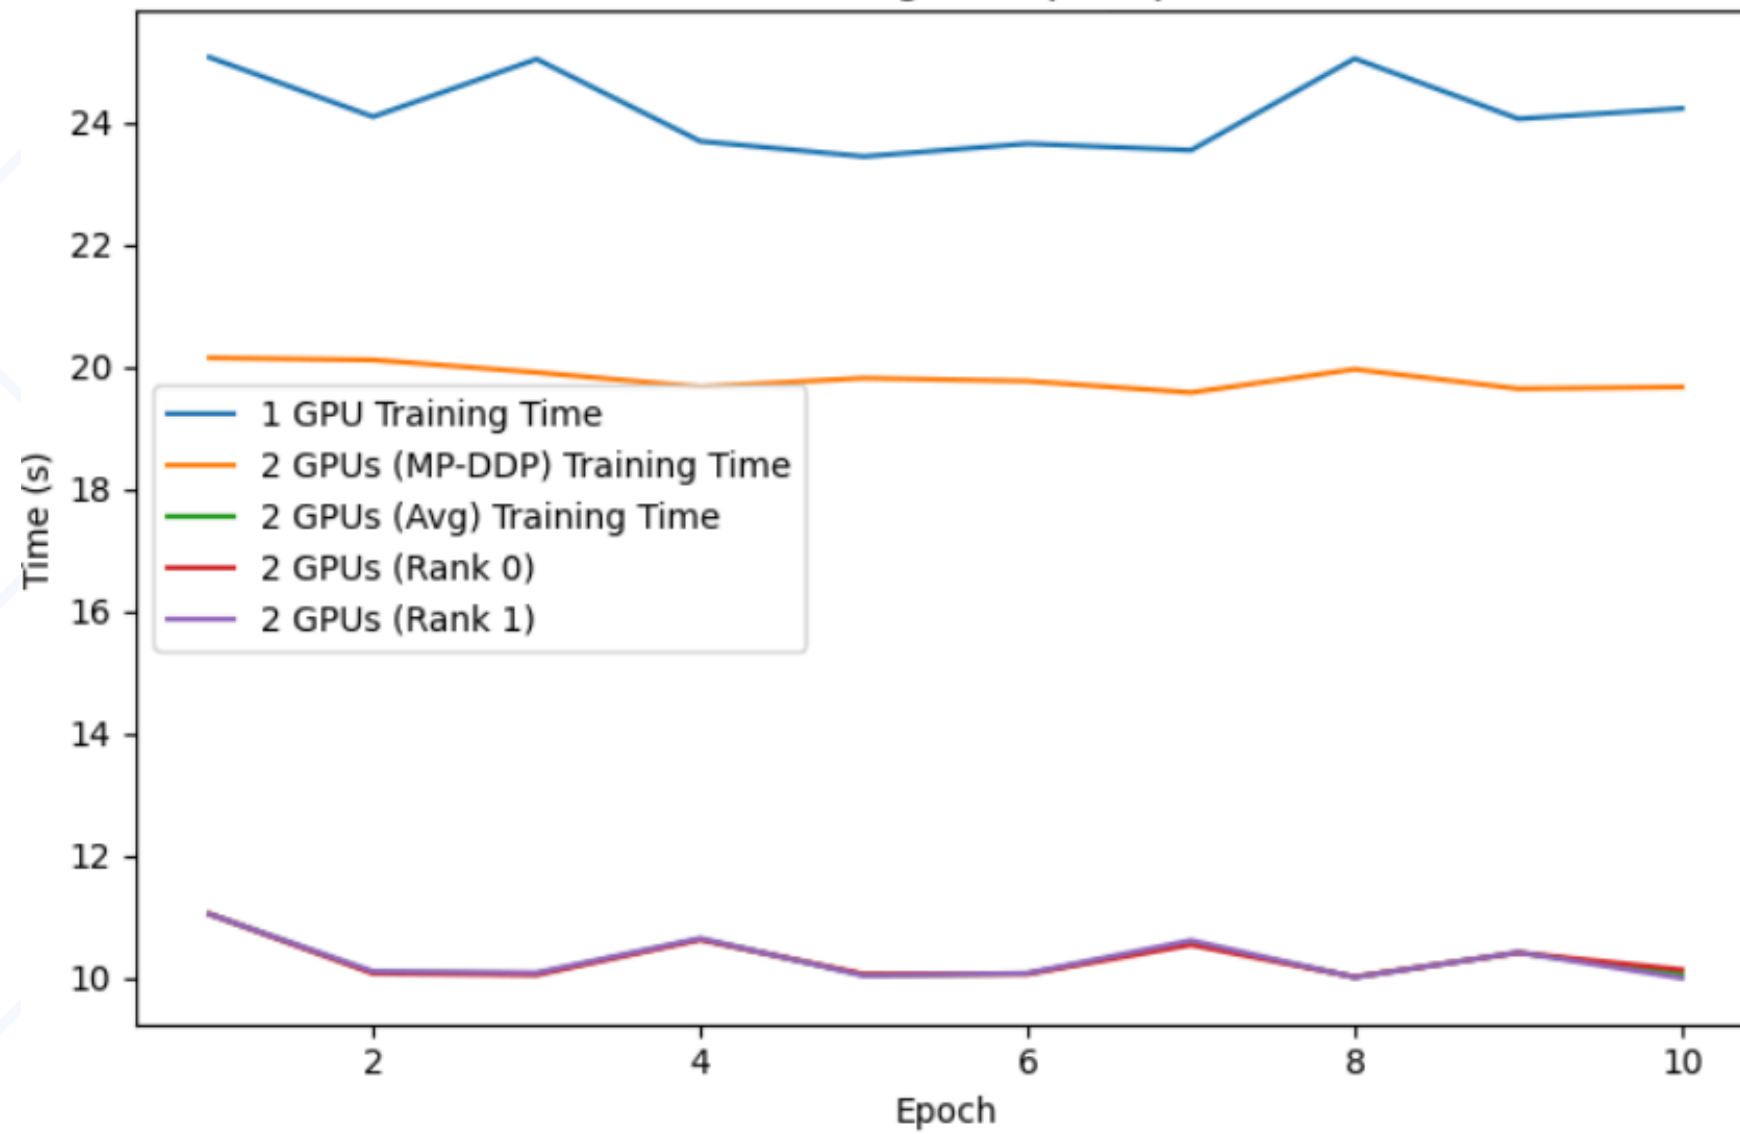


03

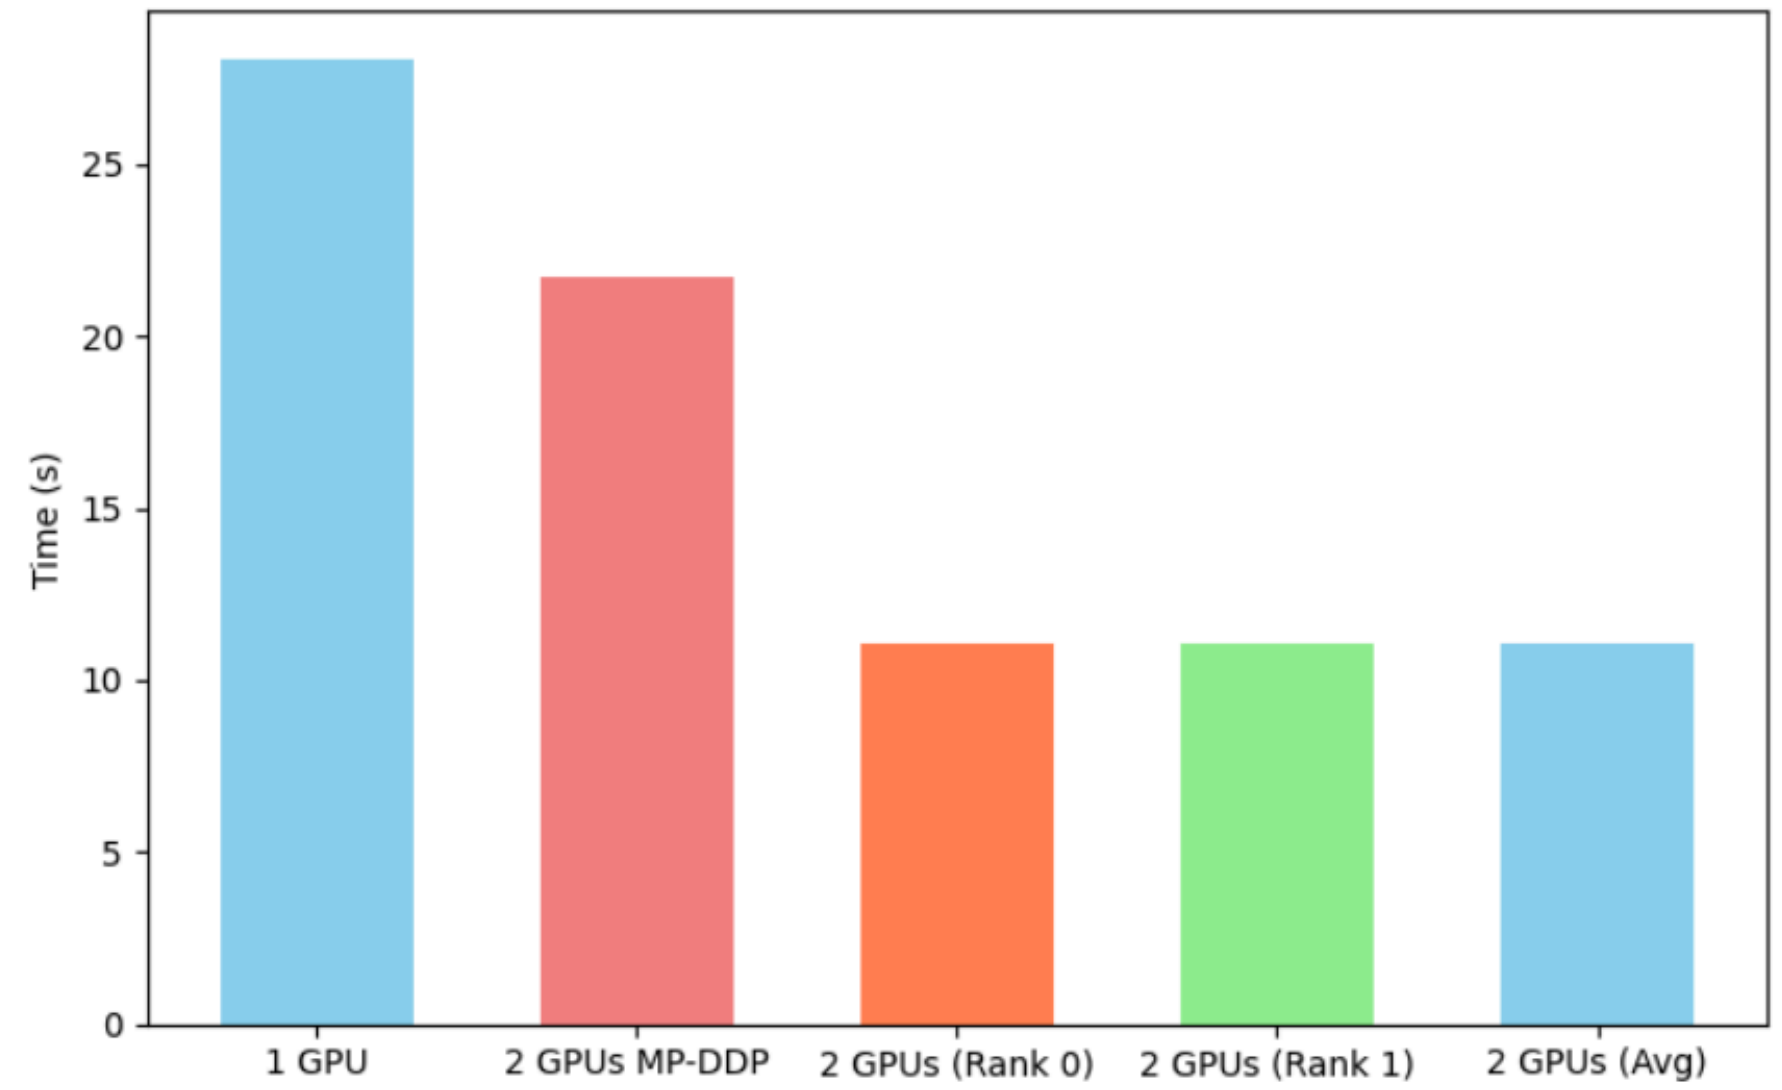
MLP

Training and Validation Time

MLP - Training Time per Epoch

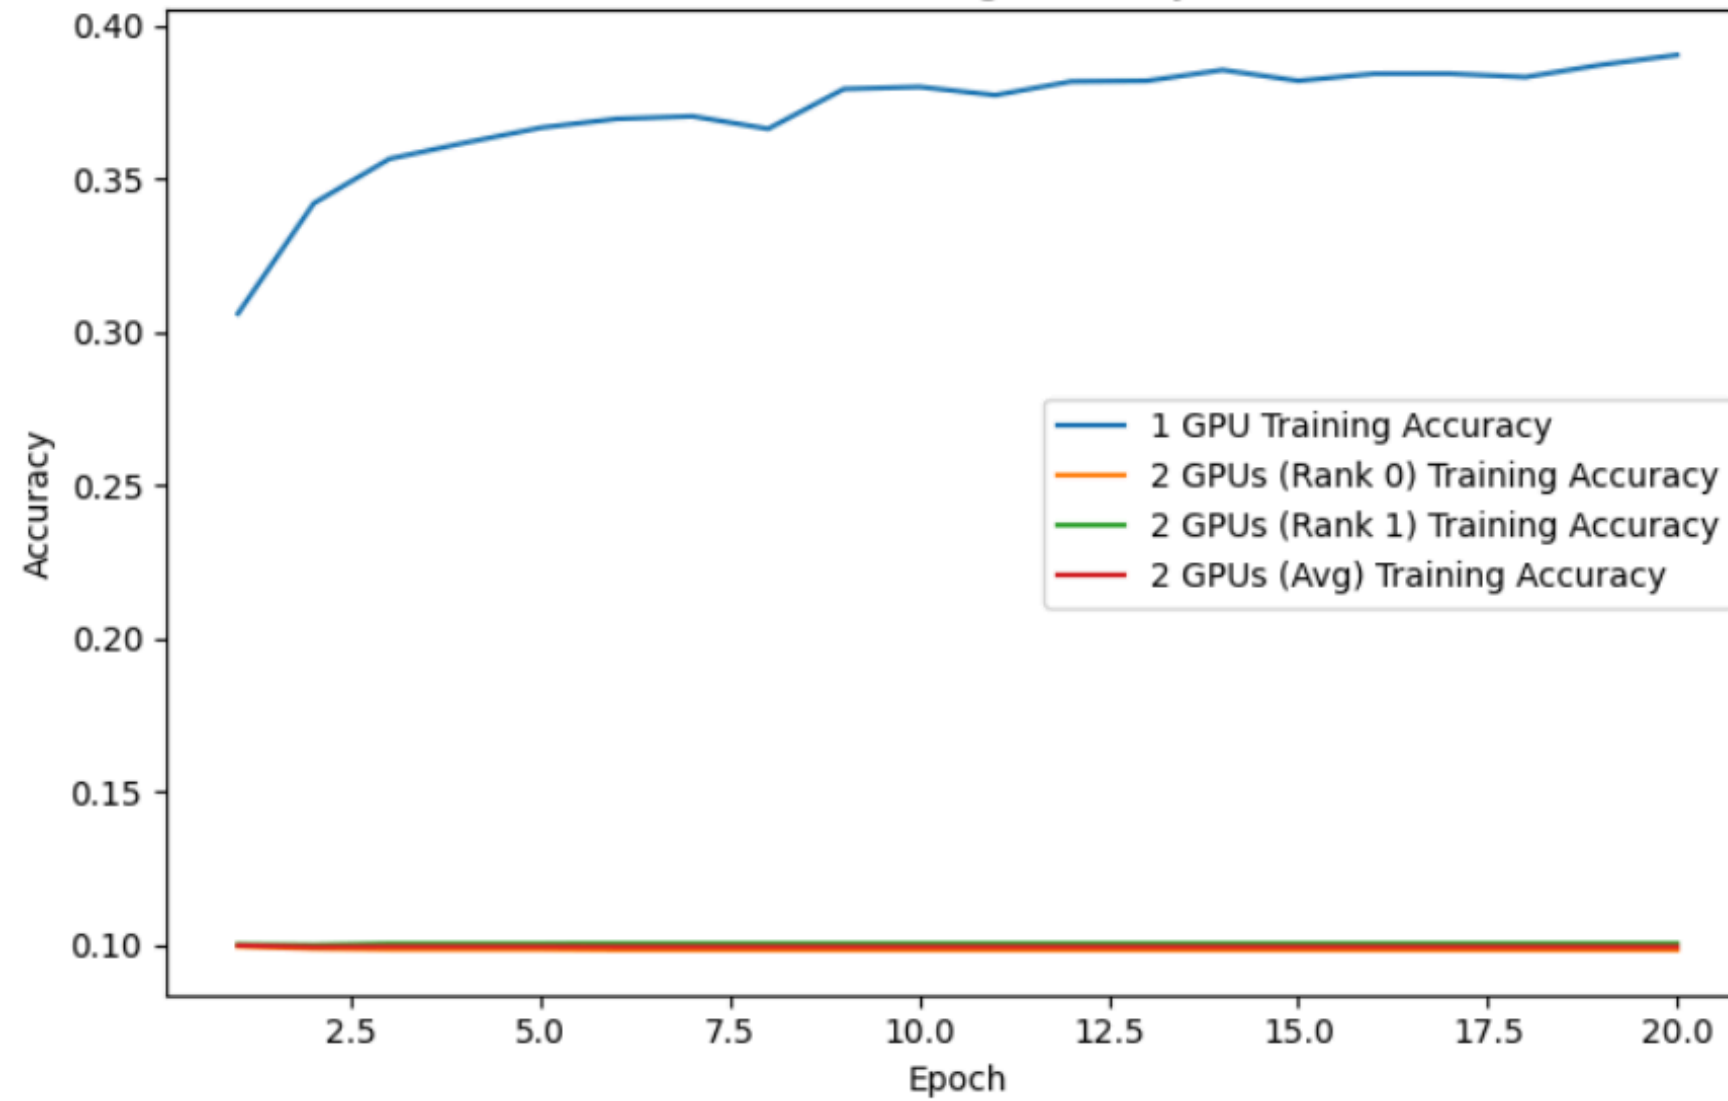


MLP - Total Validation Time




Training and Validation Accuracy

CNN - Training Accuracy

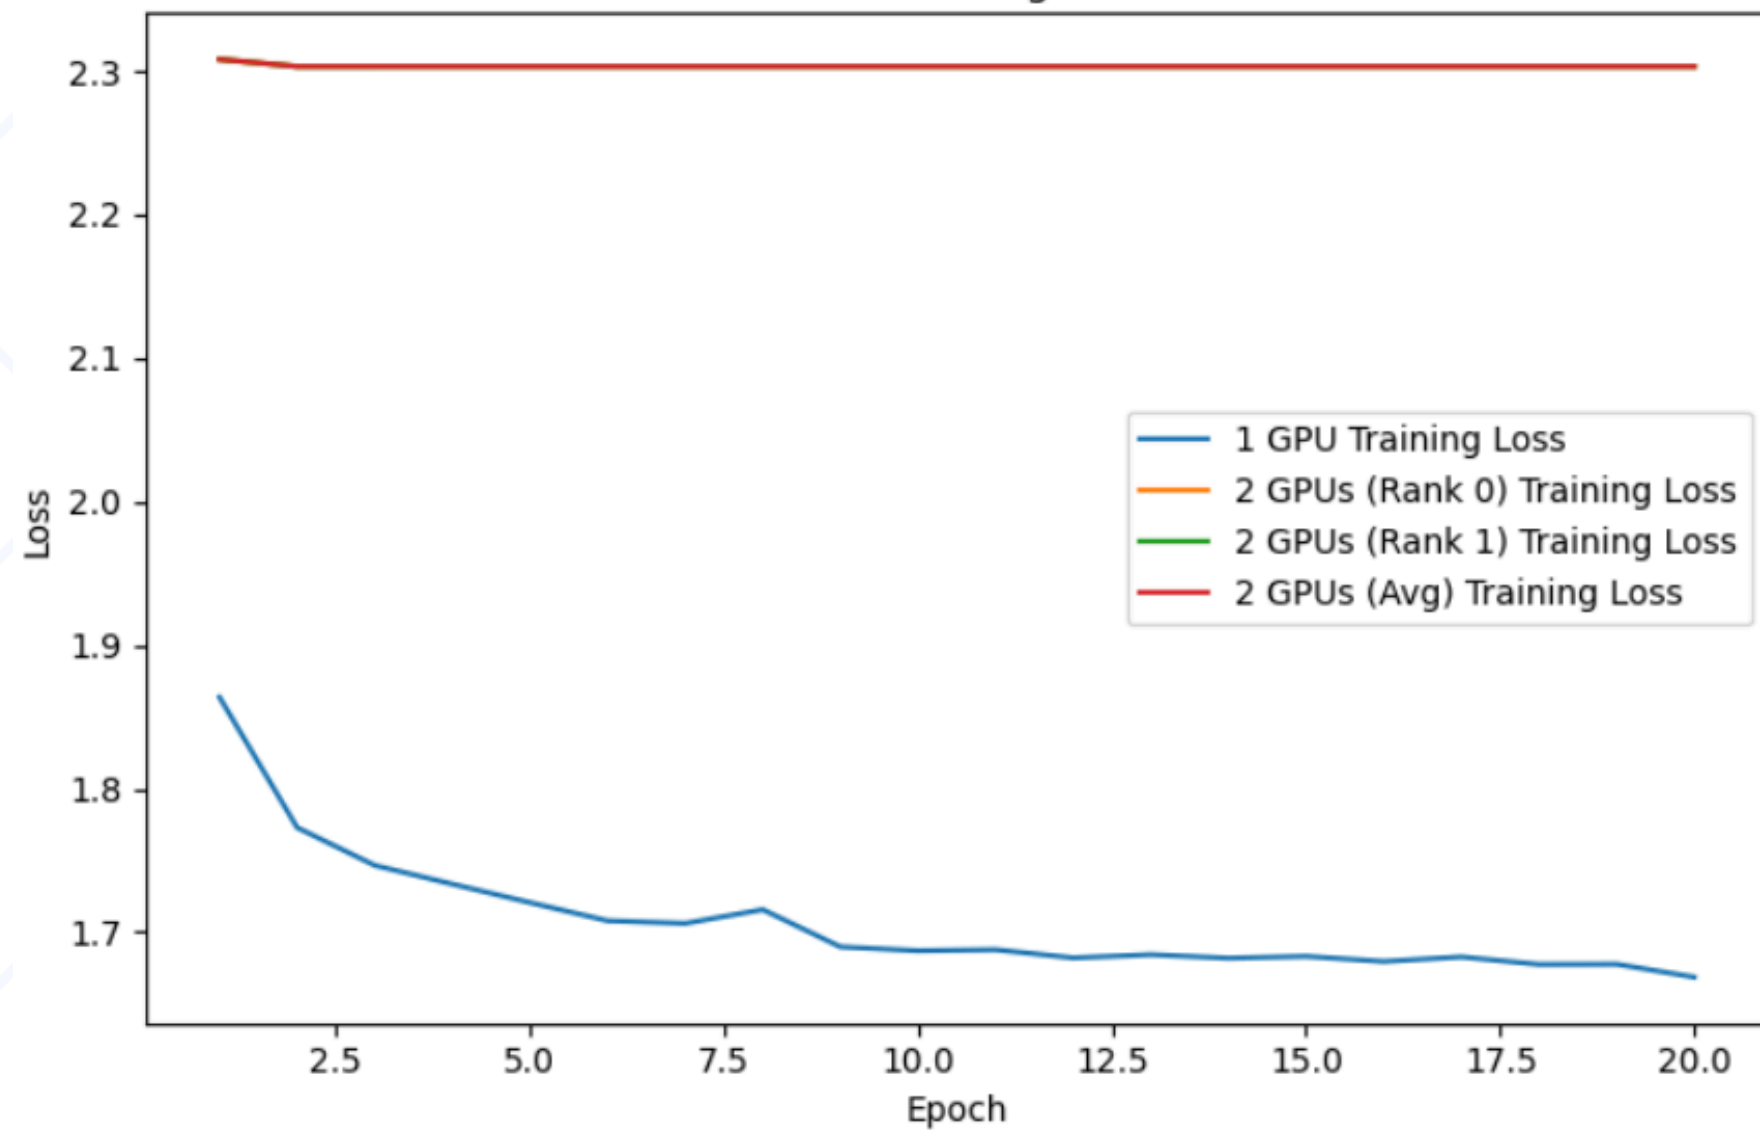


CNN - Validation Accuracy

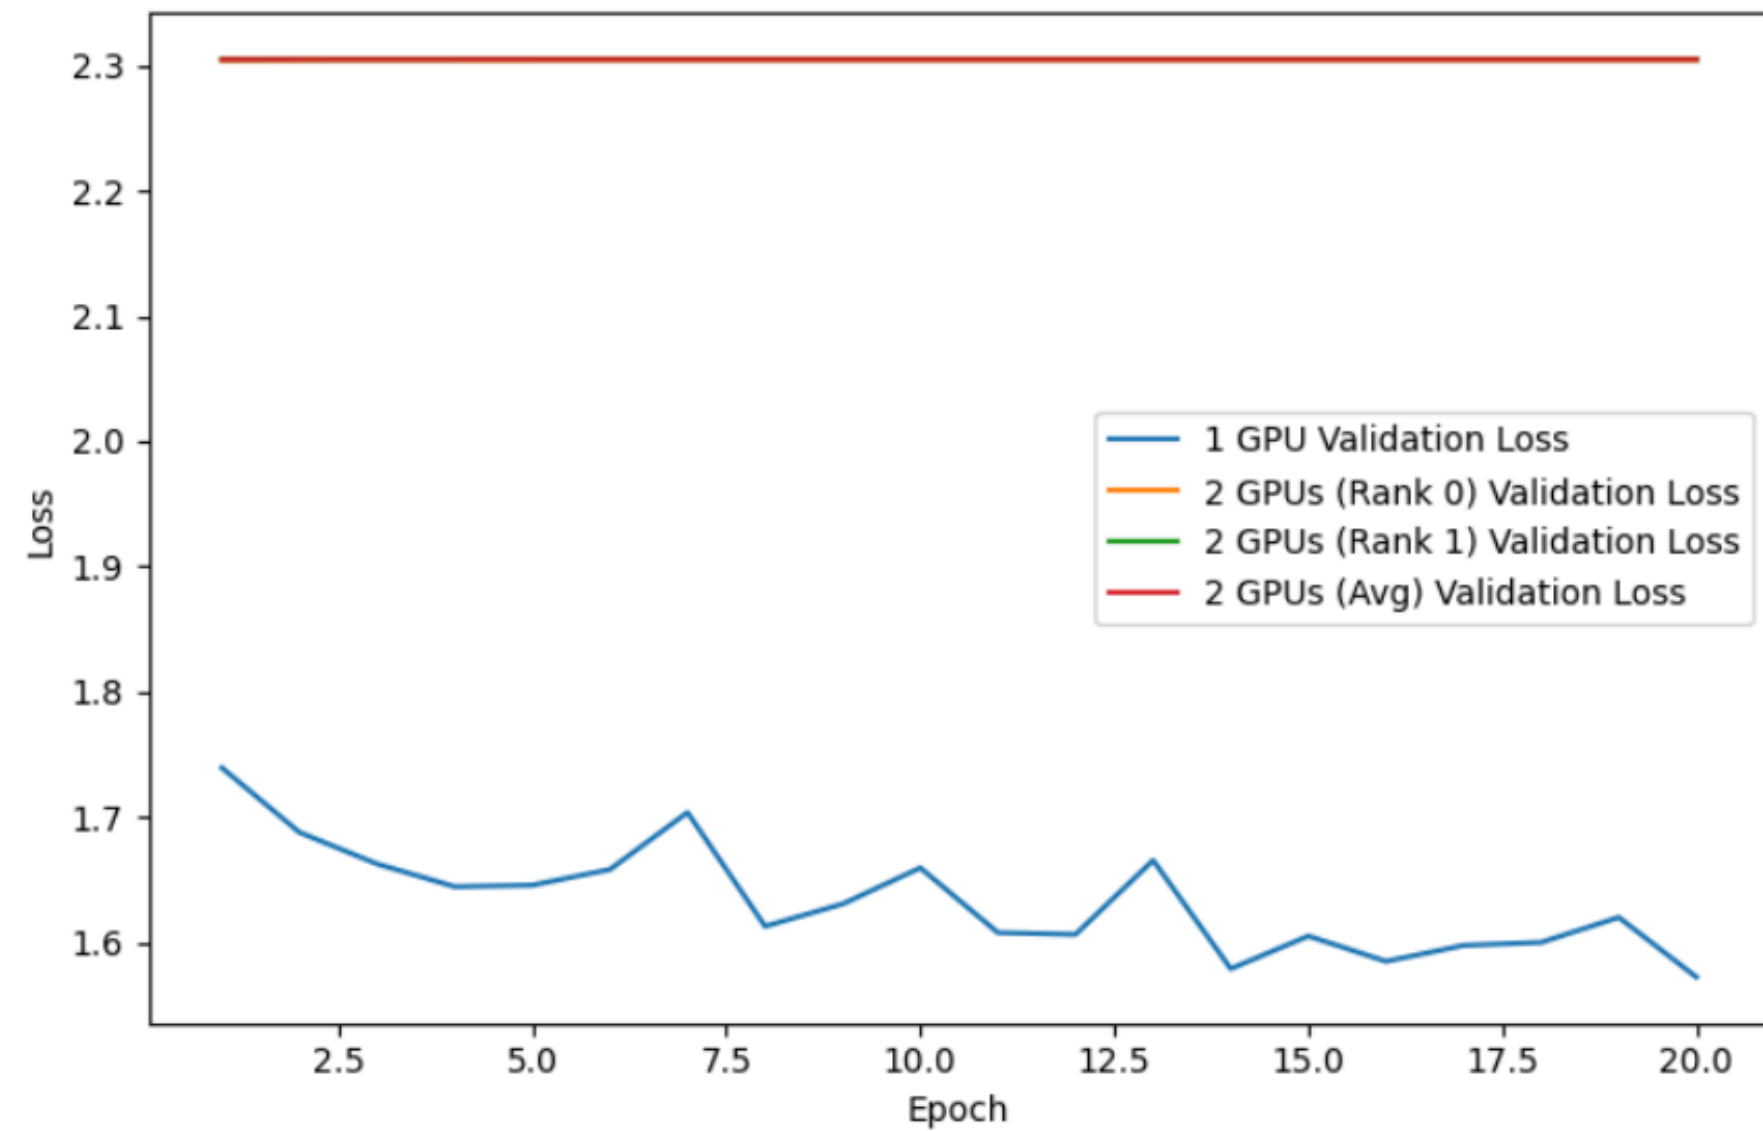


Training and Validation Loss

CNN - Training Loss

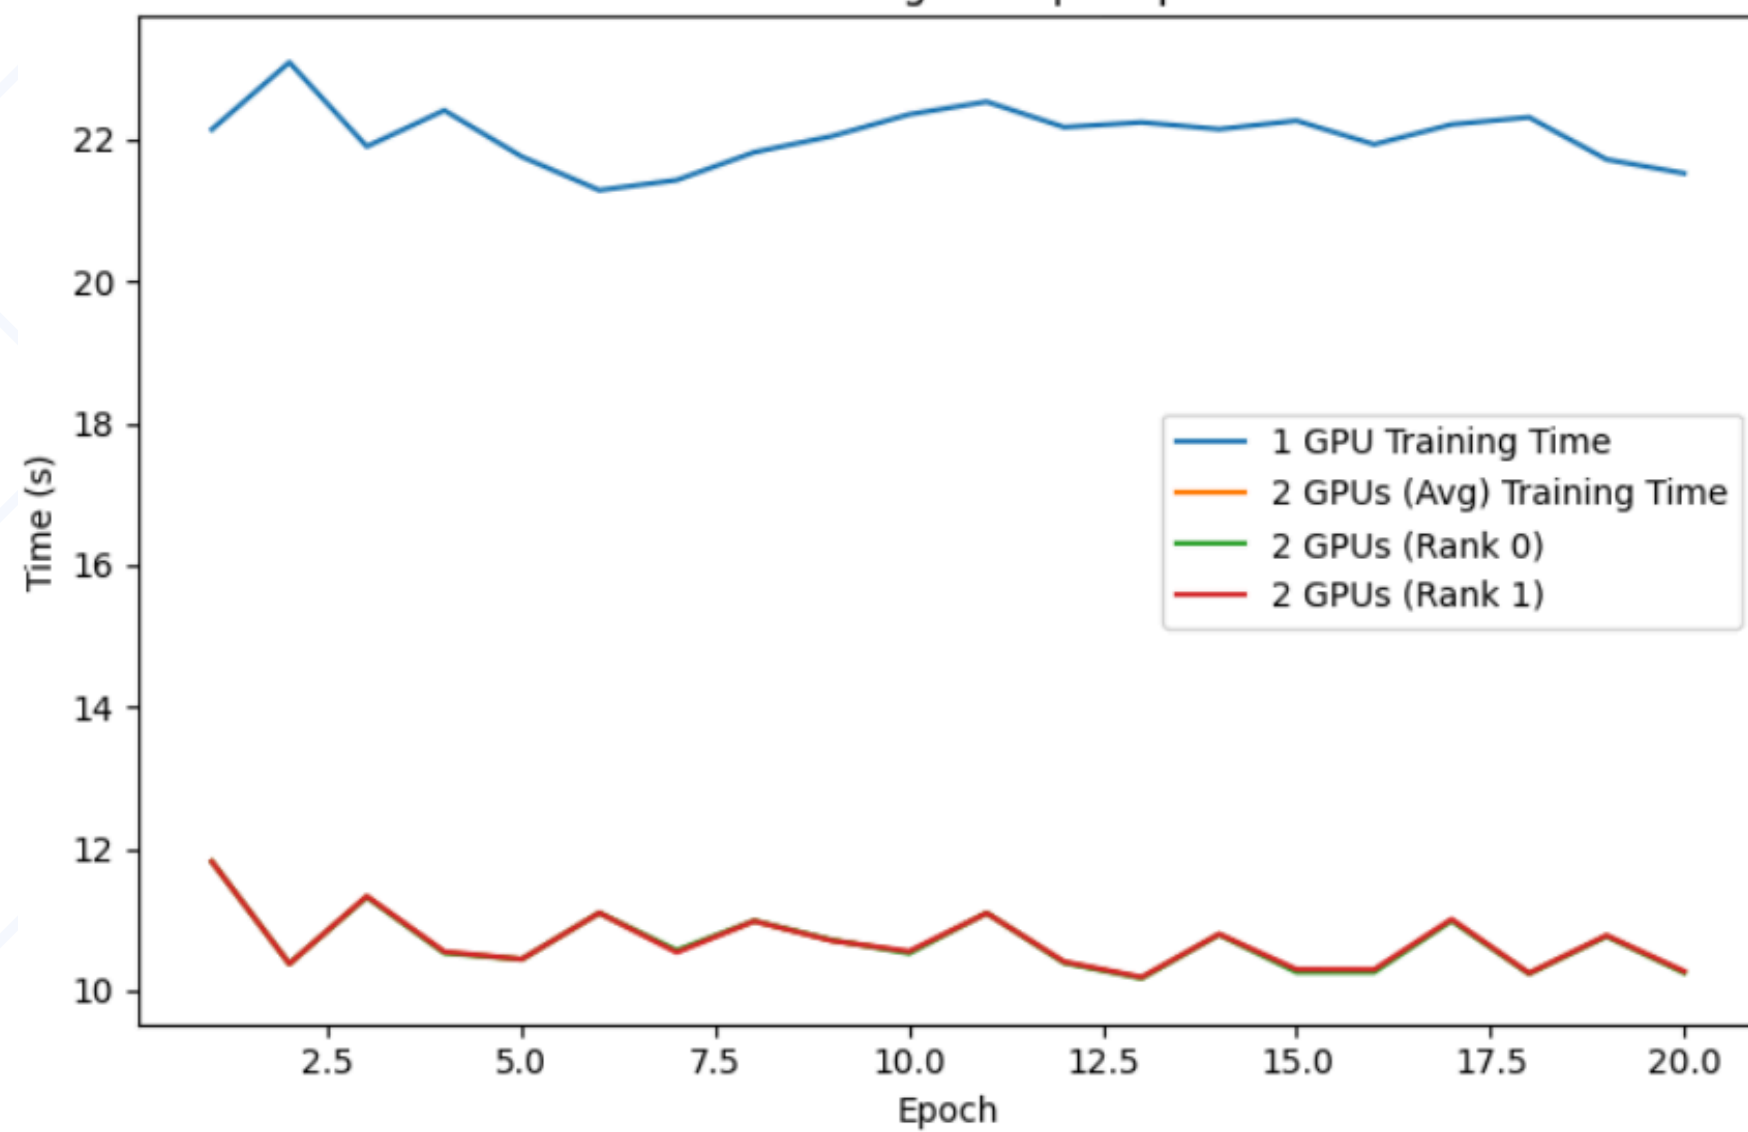


CNN - Validation Loss




Training and Validation Time

CNN - Training Time per Epoch



CNN - Total Validation Time




03


NBoW

Training and Validation Loss

NBoW - Training Loss




NBoW - Validation Loss




NBoW

Training and Validation Time

NBoW - Training Time per Epoch




NBoW - Total Validation Time




Seq 2 Seq

Training and Validation Loss

Seq2Seq - Training Loss





Seq2Seq - Validation Loss



Seq 2 Seq

Training and Validation Time




03


Convolutional Autoencoder

Training and Validation Loss

ConvAutoencoder - Training Loss





ConvAutoencoder - Validation Loss



Convolutional Autoencoder

Training and Validation Time



04

Results Analysis

Results Analysis

Metric	Training and Validation Accuracy	Training and Validation Loss	Training and Validation Time
Distributed improves the metric	VGG11, MLP (Hybrid Parallelism)	VGG11, MLP (Hybrid Parallelism)	VGG11, NBoW, Seq2Seq, ConvAutoencoder, CNN, MLP (Data Parallelism)
Distributed has no change on the metric	NBoW, Seq2Seq, ConvAutoencoder	NBoW, Seq2Seq, ConvAutoencoder	-
Distributed deteriorates the metric	CNN	CNN	-

05

Conclusion



Thank You